University of Zagreb
martin.tutek@gmail.com

Martin Tutek

https://mttk.github.io/
Google Scholar

Research Interests

Knowledge Localization and Editing, Explainability and Controllability of Language Models

Employment

Postdoctoral Researcher, University of Zagreb Feb 2025 – Present
- Substituting for prof. Jan Šnajder in 24/25 summer semester

Postdoctoral Researcher, Technion Feb 2024 – Feb 2025

Postdoctoral Researcher, UKP Lab, Technische Universität Darmstadt Sep 2022 – Dec 2023

Research Consultant, European Commission, Joint Research Centre Sep 2014 – Sep 2015

Education

Ph.D. In Computer Science, University of Zagreb 2016 – 2022

M.Sc. In Computer Science, University of Zagreb 2012 – 2014

B.Sc. In Computer Science, University of Zagreb 2009 – 2012

Publications

Selected Publications

[*] **Tutek, M.**, Hashemi Chaleshtori, F, Marasović, A, Belinkov, Y. (2025). Measuring Faithfulness of Natural Language Explanations by Unlearning Reasoning Steps. EMNLP 2025 & Interplay workshop @ COLM 2025 (*oral*).

[*] Mueller*, A., Geiger*, A., Wiegreffe*, S., …, **Tutek, M**., Zur, A, Bau, D., Belinkov, Y. (2025). MIB: A Mechanistic Interpretability Benchmark. ICML 2025 & Actionable Interpretability workshop @ ICML 2025.

[*] **Tutek, M.** & Šnajder, J. (2020). Staying True to Your Word:(How) Can Attention Become Explanation?. Representation Learning for NLP workshop @ ACL 2020.

Preprints or under review

[1] Bibhuti, A., Vashishtha, S., Naik, A., **Tutek, M.**, Aditya, S. (2025). PragWorld: A Benchmark Evaluating LLMs' local world model under Perturbations and Conversational Dynamics. Under review.

[2] Ashuach, T., Arad, D., Mueller, A., **Tutek, M.**, Belinkov, Y. (2025). CRISP: Persistent Concept Unlearning via Sparse Autoencoders. Under review.

Peer-reviewed publications

[3] **Tutek, M.**, Hashemi Chaleshtori, F, Marasović, A, Belinkov, Y. (2025). Measuring Faithfulness of Natural Language Explanations by Unlearning Reasoning Steps. EMNLP 2025 & Interplay workshop @ COLM 2025 (*oral*).

[4] Soker, Y., **Tutek, M.**, & Belinkov, Y. (2025). Predicting Success of Model Editing Through Intrinsic Features. Interplay workshop at COLM 2025.

[5] Dukić, D., Barić, A., Čuljak, M., Jukić, J., **Tutek, M.** (2025). Characterizing Linguistic Shifts in Croatian News via Diachronic Word Embeddings. Slavic NLP workshop at ACL 2025.

[6] Mueller*, A., Geiger*, A., Wiegreffe*, S., …, **Tutek, M**., Zur, A, Bau, D., Belinkov, Y. (2025). MIB: A Mechanistic Interpretability Benchmark. ICML 2025.

[7] Ashuach, T., **Tutek, M.**, & Belinkov, Y. (2024). REVS: Unlearning Sensitive Information in Language Models via Rank Editing in the Vocabulary Space. Findings of ACL 2025.

[8] Puerto, H., **Tutek, M.**, Aditya, S., Zhu, X., & Gurevych, I. (2024). Code Prompting Elicits Conditional Reasoning Abilities in Text+ Code LLMs. EMNLP 2024.

[9] Jelenić, F., Jukić, J., **Tutek, M.**, Puljiz, M., & Šnajder, J. (2024). Out-of-Distribution Detection by Leveraging Between-Layer Transformation Smoothness. ICLR 2024.

[10] Sachdeva, R., **Tutek, M.**, & Gurevych, I. (2024). CATfOOD: Counterfactual Augmented Training for Improving Out-of-Domain Performance and Calibration. EACL 2024.

[11] Jukić, J.*, **Tutek, M.***, & Šnajder, J. (2023). Easy to Decide, Hard to Agree: Reducing Disagreements Between Saliency Methods. Findings of ACL 2023 ***Equal contributions**.

[12] **Tutek, M**., & Snajder, J. (2022). Toward Practical Usage of the Attention Mechanism as a Tool for Interpretability. IEEE Access.

[13] Obadić, L., **Tutek, M.**, & Šnajder, J. (2022). NLPOP: a Dataset for Popularity Prediction of Promoted NLP Research on Twitter. In Proceedings of the 12th Workshop on Computational Approaches to Subjectivity, Sentiment & Social Media Analysis (pp. 286-292).

[14] **Tutek, M.** & Šnajder, J. (2020). Staying True to Your Word:(How) Can Attention Become Explanation?. Representation Learning for NLP Workshop @ ACL 2020.

[15] **Tutek, M.** & Šnajder, J. (2018). Iterative Recursive Attention Model for Interpretable Sequence Classification. In Proceedings of the 2018 EMNLP Workshop: Analyzing and interpreting neural networks for NLP.

[16] di Buono, M. P., Šnajder, J., Dalbelo Bašić, B., Glavaš, G., **Tutek, M.**, & Milic-Frayling, N. (2017). Predicting News Values from Headline Text and Emotions. In 2017 EMNLP Workshop on Natural Language Processing Meets Journalism (pp. 1-6).

[17] Rotim, L., **Tutek, M.**, & Šnajder, J. (2017, August). Takelab at semeval-2017 task 5: Linear aggregation of word embeddings for fine-grained sentiment analysis of financial news. In Proceedings of the 11th International Workshop on Semantic Evaluation (SemEval-2017) (pp. 866-871).

[18] di Buono, M. P., **Tutek, M.**, Šnajder, J., Glavaš, G., Bašic, B. D., & Milic-Frayling, N. (2017). Two Layers of Annotation for Representing Event Mentions in News Stories. LAW XI 2017, 82.

[19] **Tutek, M.**, Glavas, G., Šnajder, J., Milić-Frayling, N., & Dalbelo Basic, B. (2016, October). *Detecting and Ranking Conceptual Links between Texts Using a Knowledge Base*. In Proceedings of the 25th ACM International on Conference on Information and Knowledge Management (pp. 2077-2080).

[20] **Tutek, M.**, Sekulić, I., Gombar, P., Paljak, I., Čulinović, F., Boltužić, F., Karan, M., Alagić, D. and Šnajder, J. (2016). Takelab at semeval-2016 task 6: stance classification in tweets using a genetic algorithm based ensemble. In Proceedings of the 10th International Workshop on Semantic Evaluation (SemEval-2016) (pp. 464-468).

Invited Talks

From Internals to Integrity: How Insights into Transformer LMs Improve Safety and Faithfulness

- *"The Cross-Disciplinary Impact of Transformers and LLMs"* lecture series, FER, University of Zagreb — Jun 2025

Using, Improving and Calibrating (Large) Language Models

- CroAI Meetup — Mar 2024

ChatGPT: What can large pre-trained language models say about the future?

- CroAI Meetup — May 2023
- IEEE Croatia, University of Zagreb — Dec 2022

On interpretability: attention, saliency and beyond

- Technion — Jun 2022

TEACHING

209719: Introduction to Artificial Intelligence, University of Zagreb — Summer 2025

- *Lecturer (50% of course; 580 students)*

20-00-0947: Deep Learning for Natural Language Processing, TU Darmstadt — Summer 2023

- *Lecturer (50% of course; 130 students)*

252377: Deep Learning 1, University of Zagreb — Summer 2018-2021

- *Lab assignments and lectures (2 lectures)*

222925: Text Analysis and Retrieval, University of Zagreb — Autumn 2018-2021

- *Lectures (2 lectures)*

209719: Introduction to Artificial Intelligence, University of Zagreb — Summer 2016-2022

- Lab assignments

MENTORING

University of Zagreb PhD students (*with Jan Šnajder*)

- Josip Jukić (now Postdoc at University of Zagreb) — Aug 2021 – Jul 2025
  - Papers @ ACL 2023 Findings, ICLR 2024

Technion PhD Students *(with Yonatan Belinkov)*

- Tomer Ashuach (*now PhD at Technion*) — Feb 2024 – present
  - Papers @ ACL 2025 Findings, preprint

Darmstadt PhD Students *(with Iryna Gurevych)*

- Haritz Puerto — Dec 2022 – Dec 2023
  - Paper @ EMNLP 2024 Main
- Rachneet Sachdeva — Sep 2022 – Oct 2023
  - Paper @ EACL 2024 Main

SERVICE

Organizing experience:

- **BlackBoxNLP 2025** Shared task: *benchmarking new techniques for localizing circuits and causal latent variables in language models* (https://blackboxnlp.github.io/2025/task/)

Area Chair: Interpretability and Analysis of Models for NLP

- COLING 2024
- ACL Rolling Review
  Dec 2023–current
- EMNLP 2023
- ACL 2023

Conference Reviewer

- AAAI 2026
- ICLR 2025
- COLM 2024, 2025

- NeurIPS 2024, 2025
- ARR Nov *2021* – Oct 2023 *(outstanding reviewer Oct 2023)*
- EACL 2023
- EMNLP *2018* – 2022
- ACL *2018* – 2022

Journal Reviewer

- *Automatika* 2020, 2021
- *Artificial Intelligence* 2021, 2022

Workshop Reviewer

- L2M2 Workshop @ACL 2025
- Actionable Interpretability Workshop @ICML 2025
- Mechanistic Interpretability Workshop @ICML 2024
- BlackboxNLP @EMNLP 2018, EMNLP 2025

Summer School Lecturer

Intl' Summer School of Data Science in Split, practical sessions. 2016, 2017

Open-Source Software and Datasets

- **Podium:** A framework agnostic python NLP library for data loading and preprocessing
  https://github.com/TakeLab/podium
- **RNN-classifier**: A minimalistic RNN classifier sample
  https://github.com/mttk/rnn-classifier
- **pytorch**
  Contributor
  https://github.com/pytorch/pytorch
- **pytorch-text**
  Contributor, maintainer
  https://github.com/pytorch/text

References

- Yonatan Belinkov, belinkov@technion.ac.il
- Jan Šnajder, jan.snajder@fer.hr
- Ana Marasović, ana.marasovic@utah.edu